

1 Question 1: conv2d

Done as instructed; tests passed.

2 Question 2: Pararellelization

Done as instructed and tested for correctness. Marginal improvement.

3 Question 3: Tweaking hyper-parameters

I first of all changed the batch size from 512 to 128 and then went with the following architecture:

- Conv2D layer: Filter size: 6×6 , filter depth: 8, strides: 2,2, output: $12 \times 12 \times 8$
- LeakyReLU
- Conv2D layer: Filter size: 5×5 , filter depth: 16, strides: 1,1, output: $8 \times 8 \times 16$
- LeakyReLU
- Conv2D layer: Filter size: 4×4 , filter depth: 32, strides: 1,1, output: $5 \times 5 \times 32$
- LeakyReLU
- Conv2D layer: Filter size: 3×3 , filter depth: 64, strides: 2,2, output: $2 \times 2 \times 64$
- LeakyReLU
- Conv2D layer: Filter size: 2×2 , filter depth: 128, strides: 1,1, output: $1 \times 1 \times 128$
- LeakyReLU
- Flatten to fully connected linear layer: Size: 128
- LeakyReLU
- Fully connected linear layer: Size: 10
- SoftMax

Training over the same 50 epochs and overfitting a bit, a test set MSE of 0.053001 and a test set precision of 0.985377 were achieved with early stopping after 42 epochs. This is quite a bit better than the untuned test results after 50 epochs (MSE:0.107567 precision:0.966900).